

Assignment 2 – Mean-Variance Optimization Hands-on Exercise

Contents

Final Deliverable	1
Executive summary	1
Load data	2
Preliminary data manipulation (missing data particularly)	2
Stock returns and risks	3
Mean-variance Optimization – First Attempt	4
Mean-variance Optimization: Preliminary Analysis	5
Mean-Variance Optimization: Optimal Portfolios	6
Mean-Variance Optimization: Efficient Frontier	7

Final Deliverable

You may use any programming language that you prefer. In your final submission, answer the following questions:

1. How many rows of data do you delete in the price matrix? How many rows does the return matrix have? (5pt).
2. Is the covariance matrix for the 50 selected stocks positive definite? (5pt)
3. What are the mean return and variance for the minimum risk portfolio constructed from the 50 selected risky assets? (10pt)
4. Suppose you have \$1,000 and would like to invest in the minimum risk portfolio, how much do you invest in IBM? (5pt)
5. If your risk-tolerance parameter is 0.1, what are the mean return and variance of the parametric efficient portfolio that's suitable for you? (10pt)
6. Suppose you have \$1,000 and would like to invest in the parametric-efficient portfolio, how much do you invest in IBM? (5pt)
7. Plot an efficient frontier plot in the (σ, μ) -plane along with the 50 individual assets. (10pt)

Executive summary

In this exercise, we will see the following practical problems in financial analysis:

- Missing data: Two of the stocks in the dataset has some missing data. This is troublesome in many ways (for example, cannot estimate covariance matrix).
 - A simplistic solution is to truncate the dataset to a narrower observation window where all stocks have valid prices.

- Estimated covariance matrix is NOT positive definite, which may seem strange at first. There are many ways this can happen. In our case, this is because we have more stocks (about 500) than available data (about 200).
 - In practice, you may want to find ways to collect more data.
 - In this assignment, we consider mean-variance portfolio optimization for a small subset of the stocks.

Load data

The following codes show how you can clear temporary variables (e.g., some variables left from previous codes) and load data from the .csv file provided. The first ten rows and ten columns of the loaded data are displayed.

```
rm(list=ls())      # clear temporary variables
library(ggplot2)   # library for plotting

prices <- read.csv("HistoricalPrices.csv",header = T)      # read data from file
prices[1:10,1:10] # visualize the historical price data
```

##	X	ref.date	A	AAL	AAP	AAPL	ABBV	ABC	ABMD
## 1	1	2020-01-02	85.25838	28.98289	157.9097	74.44460	83.87167	83.59086	168.81
## 2	2	2020-01-03	83.88947	27.54819	157.9196	73.72084	83.07555	82.53960	166.82
## 3	3	2020-01-06	84.13747	27.21941	155.3243	74.30827	83.73117	83.74805	179.04
## 4	4	2020-01-07	84.39538	27.11978	153.4818	73.95879	83.25351	83.14873	180.35
## 5	5	2020-01-08	85.22862	27.73749	151.7186	75.14852	83.84357	83.95437	178.69
## 6	6	2020-01-09	86.56775	27.84709	151.4313	76.74473	84.48981	85.14320	183.60
## 7	7	2020-01-10	86.88517	27.21941	147.6076	76.91822	83.41273	85.51655	189.06
## 8	8	2020-01-13	86.75623	27.28915	143.8929	78.56153	82.90698	84.64212	168.10
## 9	9	2020-01-14	87.28196	27.42864	147.4987	77.50070	83.72332	87.27522	172.73
## 10	10	2020-01-15	87.90689	27.47845	148.7964	77.16856	84.72953	89.77076	177.92
##		ABT							
## 1		85.25698							
## 2		84.21763							
## 3		84.65886							
## 4		84.18821							
## 5		84.53140							
## 6		84.75692							
## 7		83.69795							
## 8		83.46262							
## 9		84.42760							
## 10		86.04251							

Preliminary data manipulation (missing data particularly)

Seeing the first ten columns and rows of the “price” matrix above, we determine that the first two columns are irrelevant so we can remove them:

Suppose A is a matrix in R, then $A[i, j]$ calls the (i, j) -th element of the matrix. Also, $A[i,]$ and $A[, j]$ calls the i -th row and the j -th column of the matrix. The minus (‘-’) sign in “prices[, -c(1,2)]” tells R to remove Column 1 & Column 2 of the “prices” matrix.

Your tasks:

1. Take a look at the columns of the “prices” matrix with titles “CARR” and “OTIS”, you should see some NAs, i.e., missing data.

- This is because CARR and OTIS joined S&P 500 in March 2020. So these two stocks' price data is available only March 20.
2. Remove just enough rows of the “prices” matrix so that all the NAs in these two stocks are removed.
 - By removing the rows, you will be removing some stock prices for other stocks too.
 - Such removal makes sure that all the stocks have the same amount of data for later analysis. This is a simplistic solution to dealing with missing data.
 - When you are done, the first stock price for “CARR” is \$11.89012.
 3. Calculate the returns matrix based on the resulting prices matrix.
 - When you are done, the first three returns for “Stock A” are, 0.01064475, -0.04649420, and 0.05728266.
 - The returns matrix should have 1 less row than

```
prices <- prices[,-c(1,2)]      # remove the first two irrelevant columns

# complete the following line
# Your codes may be longer than 1 line, but please use the variable name "prices".
prices <- prices[-c(1:53),]    # remove some rows so all stock prices are available

# complete the following line
# Your codes may be longer than 1 line, but please use the variable name "returns".
returns <- as.data.frame(diff(as.matrix(prices)))/prices[-nrow(prices),]
```

Stock returns and risks

There is nothing you need to complete in this section. This section is simply a re-run of our synchronous time activity on Jan 18, 2021. This time, the missing data is removed.

Indeed we see a positive slop, i.e., risk-reward trade-off. But again, what we see highly depends on the observation window and the observation frequency of the stock data that we use.

```
# mean return and standard deviation
mean.return <- colMeans(returns)
std.return <- apply(returns, 2, sd)
summary(lm(mean.return ~ std.return))

##
## Call:
## lm(formula = mean.return ~ std.return)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.0044625 -0.0008479  0.0000652  0.0007331  0.0067115
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.0002943  0.0001773   -1.66  0.0976 .
## std.return   0.1102628  0.0056416   19.55 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.001251 on 499 degrees of freedom
## Multiple R-squared:  0.4336, Adjusted R-squared:  0.4325
## F-statistic: 382 on 1 and 499 DF, p-value: < 2.2e-16
```

```
# Organize data for plotting
Data <- data.frame(return = mean.return, risk = std.return)

# Plot scatter plot of stock returns vs. risks & line of best fit
p1 <- ggplot(Data, aes(y=return, x=risk)) +
  geom_point(alpha = 0.5, size=2, col= "darkorange") +
  geom_smooth(formula = y~x,method = "lm", se=F, lwd=1.5, col="#67a9cf") +
  ggtitle("Linear regression model fit")

p1
```



Mean-variance Optimization – First Attempt

There is nothing you need to complete in this section.

Please pay attention to how the mean return vector and the covariance matrix are estimated. A new library is loaded to check if the estimated covariance matrix is positive definite, which is an assumption required by mean-variance optimization.

The answer turns out to be NO?!

- One reason for non-positive-definite covariance matrix is that we have more stocks (about 500) than data (about 200).
- One consequence of non-positive-definite covariance matrix is that it is not invertible.
- We need the inverse of the covariance matrix to calculate the optimal portfolios, as we will see later.

```

mu.vec <- matrix(mean.return, ncol = 1) # mean vector, make sure it is a column vector
Sigma.mat <- cov(returns)                # covariance matrix

library(matrixcalc) # a library to check if a matrix is positive definite
is.positive.definite(Sigma.mat) # the covariance matrix fails the positive definite test

## [1] FALSE

Sig.inv <- solve(Sigma.mat) # invert the covariance matrix, or report error if not possible

## Error in solve.default(Sigma.mat): system is computationally singular: reciprocal condition number =

```

Mean-variance Optimization: Preliminary Analysis

We will next consider a mean-variance optimization problem for 50 stocks instead. With less stocks to consider, the estimated covariance matrix is now positive definite. As a result, we can calculate the optimal portfolios and visualize the efficient frontier.

Your tasks:

1. Calculate the returns of the 50 stocks specified in StockNames.csv.
2. Estimate the returns and standard deviations of the 50 stocks specified in StockNames.csv.
3. Plot the risk vs. return for the 50 stocks and the resulting line-of-best-fit. *This step is done for you if you keep the variable names as suggested.*
4. Calculate the mean vector and covariance matrix for the 50 stocks. Then verify if the covariance matrix is positive definite.

```

stocks <- read.csv("StockNames.csv")[,2] # the second column contains stock names


# complete the following lines
returns <- returns[,stocks]
mean.return <- colMeans(returns)
std.return <- apply(returns, 2, sd)

Data <- data.frame(return = mean.return, risk = std.return)

ggplot(Data, aes(y=return, x=risk)) +
  geom_point(alpha = 0.5, size=2, col= "darkorange") +
  geom_smooth(formula = y~x,method = "lm", se=F, lwd=1.5, col="#67a9cf") +
  ggtitle("Linear regression model fit")

```

Linear regression model fit



```
# complete the following line
mu.vec <- matrix(mean.return, ncol = 1) # mean vector, make sure it is a column vector
Sigma.mat <- cov(returns)                # covariance matrix

is.positive.definite(Sigma.mat) # the covariance matrix fails the positive definite test

## [1] TRUE
```

As we can see, the covariance matrix passes the positive definiteness test.

Mean-Variance Optimization: Optimal Portfolios

In this Section, you will be asked to use the formulas we derived in class to perform mean-variance portfolio optimization and answer a few questions based on your calculations.

Your tasks:

1. Calculate the minimum risk portfolio as well as its mean return and standard deviation.
2. Based on your calculations, answer Questions 3 & 4 in the “Final Deliverable” section.
3. Calculate the “zero-covariance” portfolio as well as its mean return and standard deviation.
4. Based on your calculations, answer Questions 5 & 6 in the “Final Deliverable” section.

```
Sigma.inv <- solve(Sigma.mat) # calculate the inverse of the covariance matrix
ell.vec <- matrix(1, ncol = 1, nrow = 50) # define a column vector of 1's

Sigma.ell <- Sigma.inv %*% ell.vec # matrix-vector product
ell.Sigma.ell <- as.numeric(t(ell.vec) %*% Sigma.ell) # inner product of vectors
# The "as.numeric" may seem strange. But in R, dividing an 1-by-1 matrix is an
```

```
# invalid calculation. So we need to translate the 1-by-1 matrix into a number first.
```

```
#####
```

```
# complete the following lines #
```

```
#####
```

```
x.m <- Sigma.ell / ell.Sigma.ell           # min-risk portfolio  
mu.m <- as.numeric(t(x.m) %*% mu.vec)      # mean return of min-risk portfolio  
sigma2.m <- as.numeric(t(x.m) %*% Sigma.mat %*% x.m) # variance of min-risk portfolio  
mu.m
```

```
## [1] 0.0007479783
```

```
sigma2.m
```

```
## [1] 7.232177e-05
```

```
x.m["IBM",]
```

```
##          IBM
```

```
## -0.06548979
```

```
Sigma.mu <- Sigma.inv %*% mu.vec  
ell.Sigma.mu <- as.numeric(t(ell.vec) %*% Sigma.mu)
```

```
x.z <- Sigma.mu - (ell.Sigma.mu) * x.m  
mu.z <- as.numeric(t(x.z) %*% mu.vec)  
sigma2.z <- as.numeric(t(x.z) %*% Sigma.mat %*% x.z)
```

```
tau <- 0.1  
x.tau <- x.m + tau*x.z  
mu.tau <- as.numeric(t(x.tau) %*% mu.vec)  
sigma2.tau <- as.numeric(t(x.tau) %*% Sigma.mat %*% x.tau)  
mu.tau
```

```
## [1] 0.01402798
```

```
sigma2.tau
```

```
## [1] 0.001400322
```

```
x.tau["IBM",]
```

```
##          IBM
```

```
## -0.8614289
```

Mean-Variance Optimization: Efficient Frontier

In this Section, you will be asked to use the optimal portfolios calculated above to plot the efficient frontier.

Your tasks:

1. Based on the set of risk-tolerance parameters provided, calculate the corresponding parametric efficient portfolios' mean returns and variances.
 - Hint: You can, but do not have to calculate the optimal portfolios first.
2. Plot the efficient frontier in the (σ, μ) -plane along with the 50 individual assets. Also indicate the parametric-efficient portfolio for $\tau = 0.1$ in the plot.

```
tau <- seq(0, 0.15, length.out = 100)
```

```
#####
# complete the following lines #
#####

mu.p <- mu.m + tau * mu.z
sigma2.p <- sigma2.m + tau^2 * sigma2.z

Data2 <- data.frame(mu = mu.p, sigma = sqrt(sigma2.p))
Data3 <- data.frame(mu = mu.tau, sigma = sqrt(sigma2.tau))

ggplot(Data, aes(y=return, x=risk)) +
  geom_point(alpha = 0.5, size=2, col= "darkorange") +
  geom_line(data = Data2, aes(x=sigma, y=mu), color = "#67a9cf", lwd = 1.5) +
  geom_point(data = Data3, aes(x=sigma, y=mu), color = "red", size = 5, shape=13) +
  ggtitle("Individual assets and risky efficient frontier")
```

